

---

# Feed-Forward Networks with Attention Can Solve Some Long-Term Memory Problems

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

Recently, recurrent neural networks (RNNs) have been augmented with “attention” mechanisms which compute a fixed-length representation of entire sequences. We propose a simplified model of attention which is applicable to feed-forward neural networks and demonstrate that it can solve some long-term memory problems (specifically, those where temporal order doesn’t matter). In fact, we show empirically that our model can solve these problems for sequence lengths which are both longer and more widely varying than has been shown for RNNs.

## 1 Models for Sequential Data

Many problems in machine learning are best formulated using sequential data, i.e. data where a given observation may be dependent on previous observations. Such problems can be coarsely classified as sequence transduction (producing a new sequence given an input sequence), sequence classification (producing a single label or vector from an entire sequence), or sequence generation (producing a sequence from no input) tasks. Appropriate models for these tasks must be able to capture temporal dependencies in sequences, potentially of arbitrary length.

### 1.1 Recurrent Neural Networks

One such class of models are recurrent neural networks (RNNs). Recent advances in optimization techniques [], initialization methods [], and model architectures [] have facilitated the success of RNNs in a wide variety of fields, including machine translation [], speech recognition [], handwriting recognition [], ... This variety of applications demonstrates their applicability to most machine learning tasks involving sequences. In their simplest form, RNNs can be considered as a function  $f$  whose output at time  $t$  depends on input  $x_t$ , the previous state  $h_{t-1}$ , and the model’s learnable parameters  $\theta$ :

$$h_t = f(x_t, h_{t-1}, \theta)$$

In the supervised setting, the parameters  $\theta$  are optimized with respect to a loss function which measures the performance of  $f$ . A common approach is to use backpropagation through time [], which “unrolls” the RNN over time steps to compute the gradient of  $\theta$  with respect to the loss.

Despite the success of RNNs, their recursive nature makes them hard to both optimize and parallelize. Because the same function  $f$  is applied repeatedly over time, the gradient with respect to  $\theta$  can easily explode or vanish []. The use of gating architectures [], sophisticated optimization techniques [], gradient clipping [], and/or careful initialization [] can help mitigate this issue. However, these approaches do not solve this problem, and as a result RNNs are in practice typically only applied in tasks where sequential dependencies span at most hundreds of time steps []. Very long sequences can also make training computationally inefficient due to the fact that RNNs must be evaluated sequentially and cannot be fully parallelized.

054	Attention
055	Simplified attention
056	
057	
058	<b>2 Toy Long-Term Memory Problems</b>
059	
060	Different tasks
061	Previous results
062	
063	<b>2.1 Fixed-Length Experiment</b>
064	
065	<b>2.2 Variable-length Experiment</b>
066	
067	<b>3 Limitations</b>
068	
069	
070	
071	
072	
073	
074	
075	
076	
077	
078	
079	
080	
081	
082	
083	
084	
085	
086	
087	
088	
089	
090	
091	
092	
093	
094	
095	
096	
097	
098	
099	
100	
101	
102	
103	
104	
105	
106	
107	